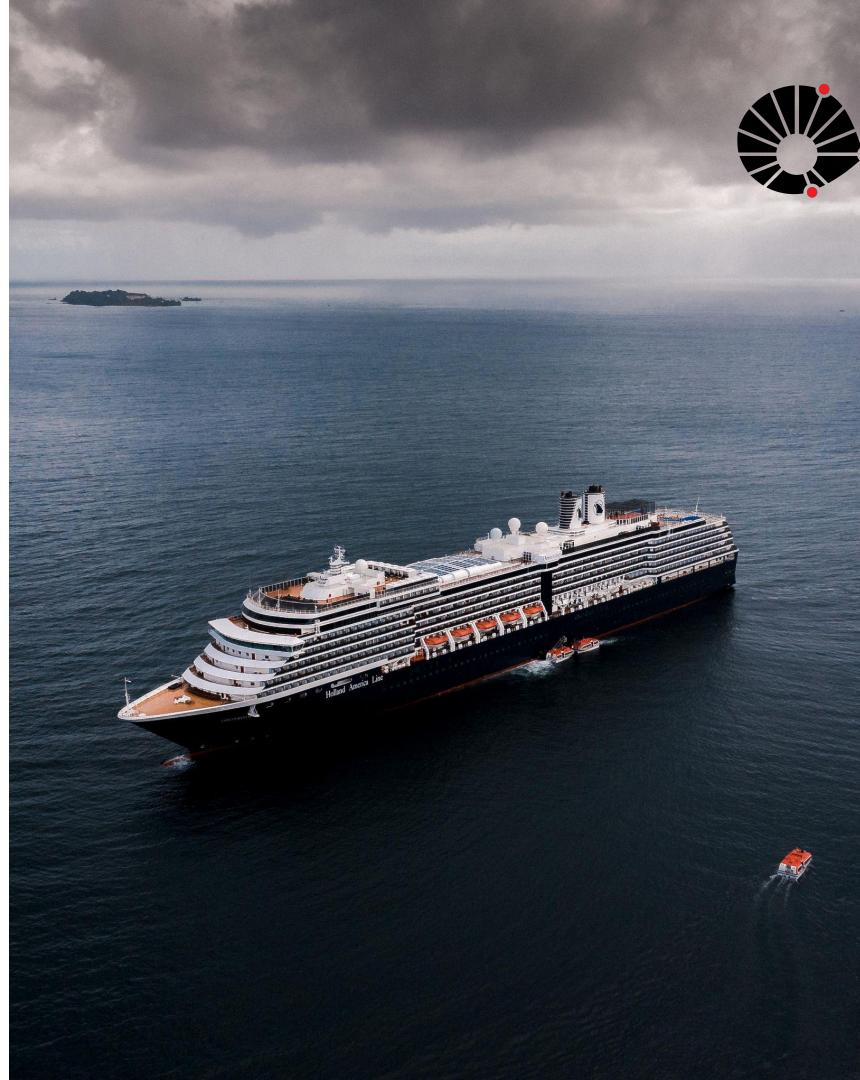


# MC714 | Avaliação 1

## titanic

Gabriel Borges, **RA 197458**  
Paulo Pacitti, **RA 185447**



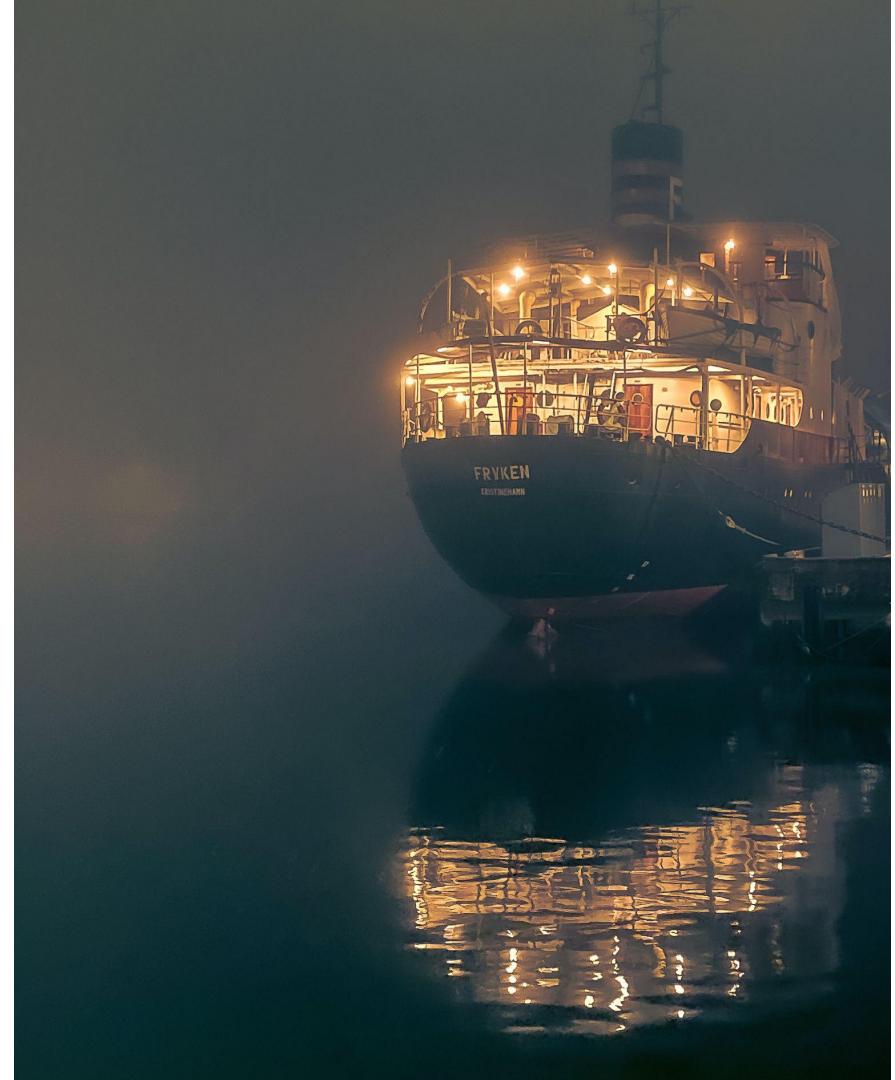
# O problema



- Implementar um cenário real de utilização de **computação em nuvem** de um **servidor web**.
- O objetivo é criar um **conjunto de máquinas virtuais que executam servidores web** que atendem clientes sem tornarem-se um gargalo.
- Para isso, o conjunto de máquinas virtuais deve ser atrelado a um **balanceador de carga** (*Load Balancer*).

# Solução Planejada

- **Aplicação:** aplicação React envelopada com Docker.
- **Infraestrutura:** Kubernetes utilizando a *Google Kubernetes Engine (GKE)* do *Google Cloud Platform (GCP)*.
- **Ferramental:** infraestrutura como código com Terraform, build e registry da imagem Docker da aplicação com Github Actions.

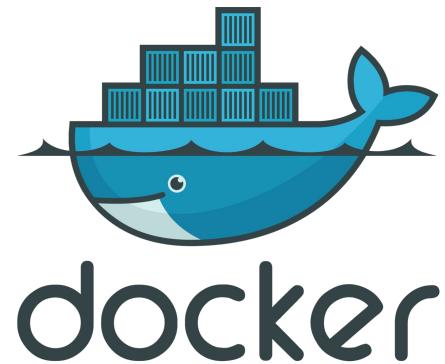
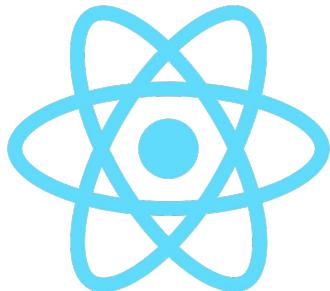


A large cargo ship, the "FUSILADA", is shown sailing on a dark blue ocean towards the left. The ship is heavily loaded with colorful shipping containers stacked high on its deck. It leaves a white wake behind it. In the background, a massive, steep hillside covered in a dense forest rises against a darkening sky. The overall atmosphere is moody and industrial.

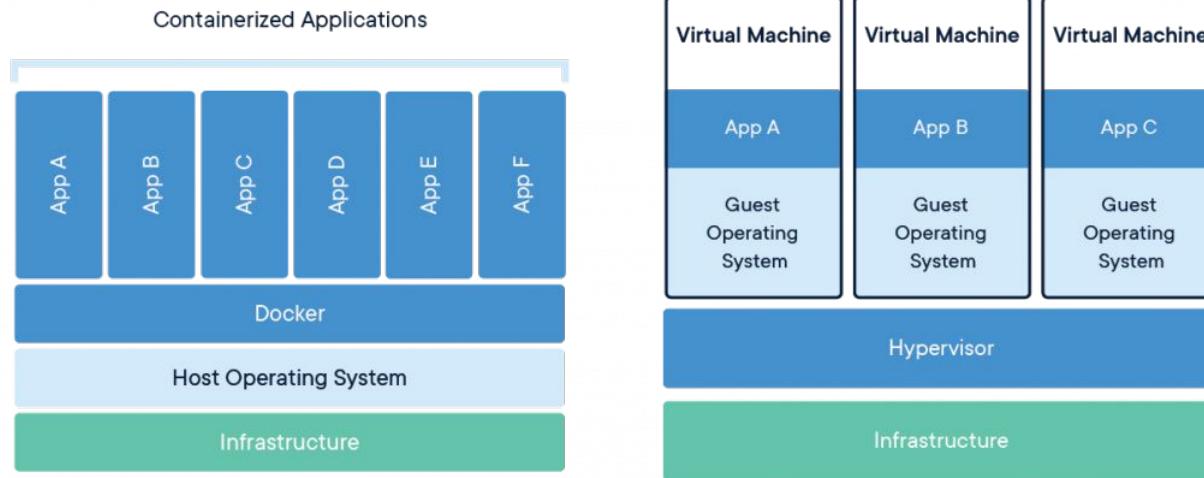
# Implementação

# Aplicação: frontend

A ideia inicial foi criar uma aplicação Web que mostra o IP de seu host. Para isso, utilizamos de **React** (*framework JavaScript*) para criar a página Web. Após isso, fazemos o *build* do site estático e envelopamos em um *container Docker* com um servidor HTTP **nginx**.



# Aplicação: Docker



Docker is a set of platform as a service (PaaS) products that use **OS-level virtualization** to deliver software in packages called ***containers***. **Containers** are **isolated from one another** and bundle their own software, libraries and configuration files; they can communicate with each other through well-defined channels. **Because all of the containers share the services of a single operating system kernel, they use fewer resources than virtual machines.**

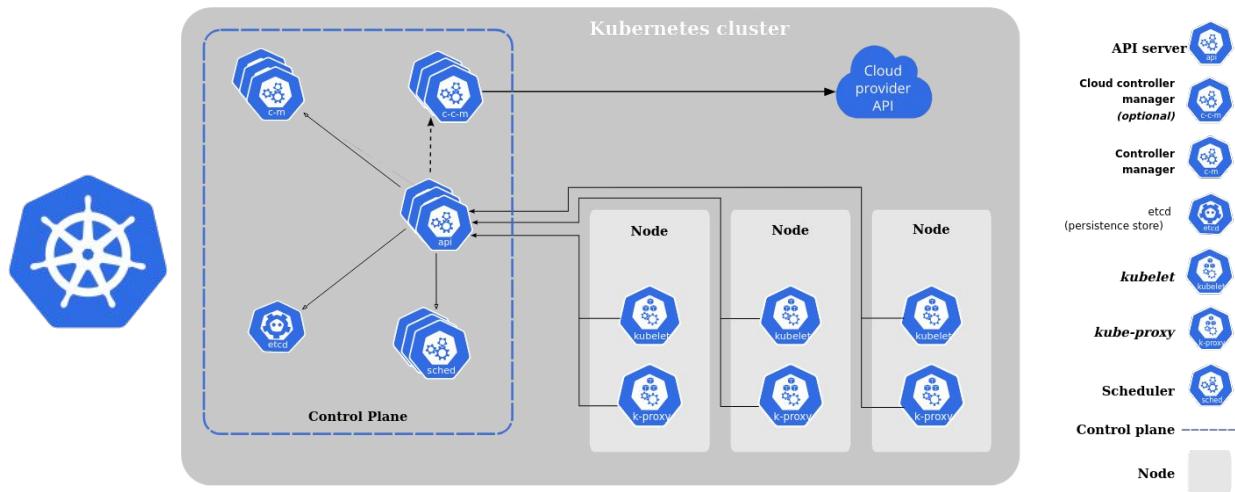
# Aplicação: Docker

```
# build environment
FROM node:13-alpine as build
WORKDIR /app
ENV PATH /app/node_modules/.bin:$PATH
COPY package.json /app/package.json
RUN npm install
RUN npm install react-scripts@3.0.1 -g
COPY . /app
RUN npm run build

# production environment
FROM nginx:alpine
COPY --from=build /app/build /usr/share/nginx/html
```

# Infraestrutura: k8s

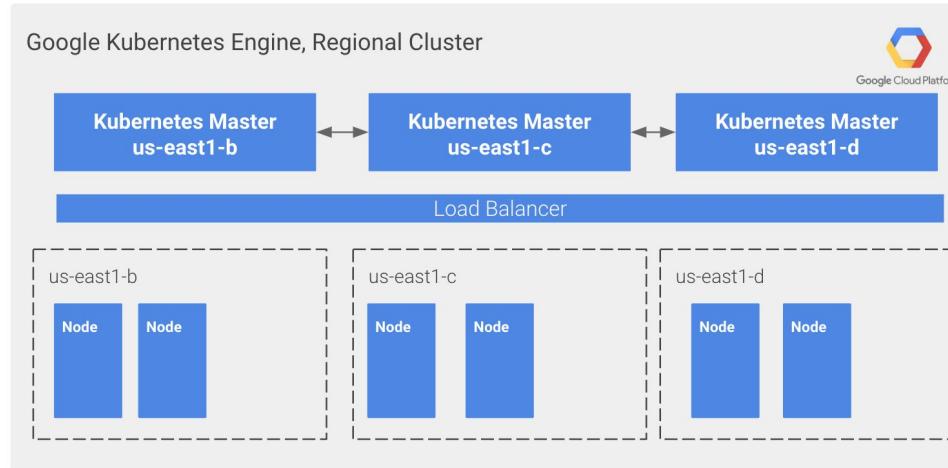
Kubernetes (k8s) is an **open-source container-orchestration system for automating computer application deployment, scaling, and management**. It was originally designed by Google and is now maintained by the Cloud Native Computing Foundation. It aims to provide a "**platform for automating deployment, scaling, and operations of database management systems**". It works with a range of container tools and **runs containers in a cluster**, often with images built using **Docker**.



# Infraestrutura: GKE

**Google Kubernetes Engine (GKE)** provides a **managed environment for deploying, managing, and scaling your containerized applications using Google infrastructure**. The GKE environment consists of multiple machines (specifically, **Compute Engine instances**) grouped together to form a cluster.

**Features:** *Identity and access management, Hybrid networking, Security and compliance, Integrated logging and monitoring, Cluster options, Auto upgrade, Auto repair...*



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
  namespace: default
spec:
  progressDeadlineSeconds: 600
  replicas: 3
  selector:
    matchLabels:
      app: frontend
  strategy:
    rollingUpdate:
      maxSurge: 33%
      maxUnavailable: 33%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: frontend
    spec:
      containers:
        - image: ghcr.io/paulopacitti/titanic:main
          imagePullPolicy: IfNotPresent
        name: frontend
        resources:
          requests:
            memory: "2Mi"
            cpu: "1m"
          limits:
            memory: "256Mi"
            cpu: "250m"
        ports:
          - containerPort: 80
            protocol: TCP
```

**deployment.yaml**

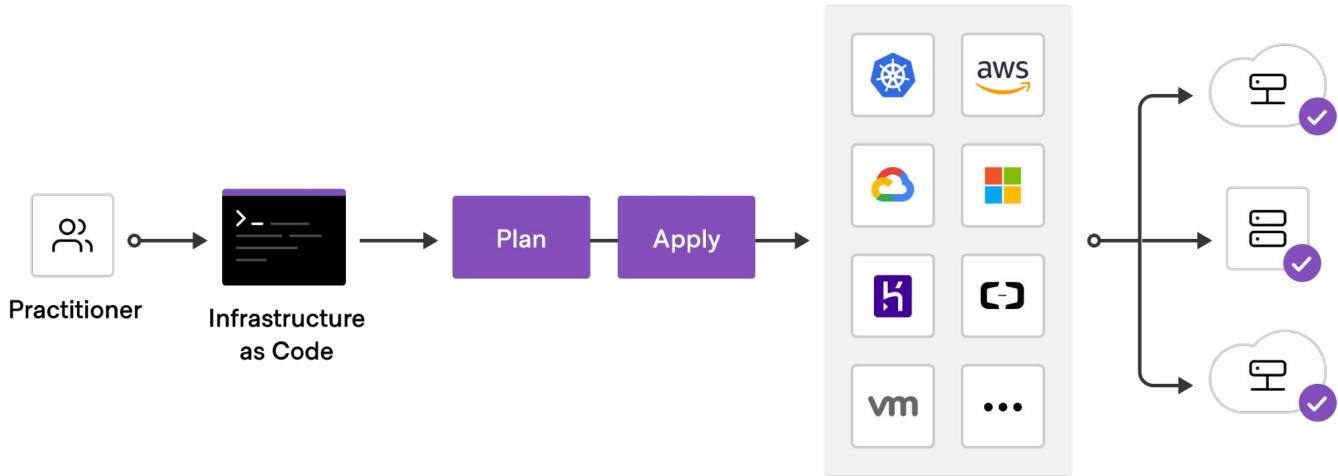
```
apiVersion: v1
kind: List
items:
  - apiVersion: autoscaling/v1
    kind: HorizontalPodAutoscaler
    metadata:
      name: frontend-hpa
      namespace: default
    spec:
      maxReplicas: 6
      minReplicas: 3
      scaleTargetRef:
        apiVersion: apps/v1
        kind: Deployment
        name: frontend
      targetCPUUtilizationPercentage: 80
```

**hpa.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: frontend-service
  namespace: default
spec:
  externalTrafficPolicy: Cluster
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    app: frontend
  sessionAffinity: None
  type: LoadBalancer
```

**service.yaml**

# Infraestrutura: Terraform



Terraform is an **open-source infrastructure as code** software tool created by HashiCorp. **Users define and provide data center infrastructure using a declarative configuration language** known as HashiCorp Configuration Language (HCL),

### versions.tf

```
terraform {  
  required_providers {  
    google = {  
      source  = "hashicorp/google"  
      version = "3.52.0"  
    }  
  }  
  
  required_version = "~> 0.14"  
}
```

```
project_id = "mc714-titanic"  
region     = "us-central1"
```

### terraform.tfvars

### outputs.tf

```
output "region" {  
  value      = var.region  
  description = "GCloud Region"  
}  
  
output "project_id" {  
  value      = var.project_id  
  description = "GCloud Project ID"  
}  
  
output "kubernetes_cluster_name" {  
  value      = google_container_cluster.primary.name  
  description = "GKE Cluster Name"  
}  
  
output "kubernetes_cluster_host" {  
  value      = google_container_cluster.primary.endpoint  
  description = "GKE Cluster Host"  
}
```

# vpc.hcl

- A virtual private cloud (VPC) is **an on-demand configurable pool of shared resources allocated within a public cloud environment**, providing a certain level of isolation between the different organizations (denoted as users hereafter) using the resources.

```
variable "project_id" {
  description = "project id"
}

variable "region" {
  description = "region"
}

provider "google" {
  project = var.project_id
  region  = var.region
}

# VPC
resource "google_compute_network" "vpc" {
  name          = "${var.project_id}-vpc"
  auto_create_subnetworks = "false"
}

# Subnet
resource "google_compute_subnetwork" "subnet" {
  name          = "${var.project_id}-subnet"
  region        = var.region
  network       = google_compute_network.vpc.name
  ip_cidr_range = "10.10.0.0/24"
}
```

```
● ● ●
variable "gke_username" {
  default    = ""
  description = "gke username"
}

variable "gke_password" {
  default    = ""
  description = "gke password"
}

variable "gke_num_nodes" {
  default    = 1
  description = "number of gke nodes"
}

# GKE cluster
resource "google_container_cluster" "primary" {
  name      = "${var.project_id}-gke"
  location  = var.region

  remove_default_node_pool = true
  initial_node_count       = 1

  network    = google_compute_network.vpc.name
  subnetwork = google_compute_subnetwork.subnet.name
}
```

```
# Separately Managed Node Pool
resource "google_container_node_pool" "primary_nodes" {
  name      = "${google_container_cluster.primary.name}-node-pool"
  location  = var.region
  cluster   = google_container_cluster.primary.name
  node_count = var.gke_num_nodes

  node_config {
    oauth_scopes = [
      "https://www.googleapis.com/auth/logging.write",
      "https://www.googleapis.com/auth/monitoring",
    ]
  }

  labels = {
    env = var.project_id
  }

  preemptible = true
  machine_type = "n1-standard-1"
  tags         = ["gke-node", "${var.project_id}-gke"]
  metadata = {
    disable-legacy-endpoints = "true"
  }
}
```

**gke.hcl**

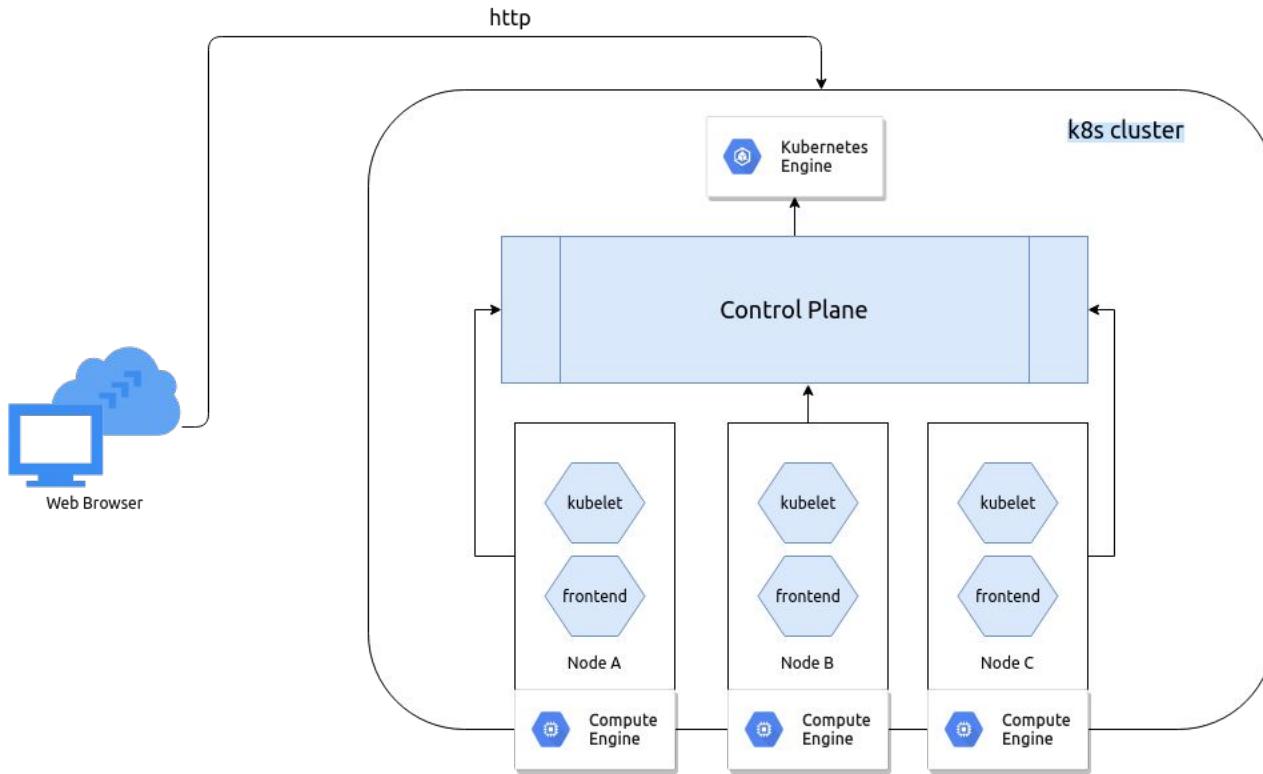
# deploy\_k8s.sh

```
#!/bin/bash
cd ../infra
CLUSTER_NAME=$(terraform output kubernetes_cluster_name)
REGION=$(terraform output region)

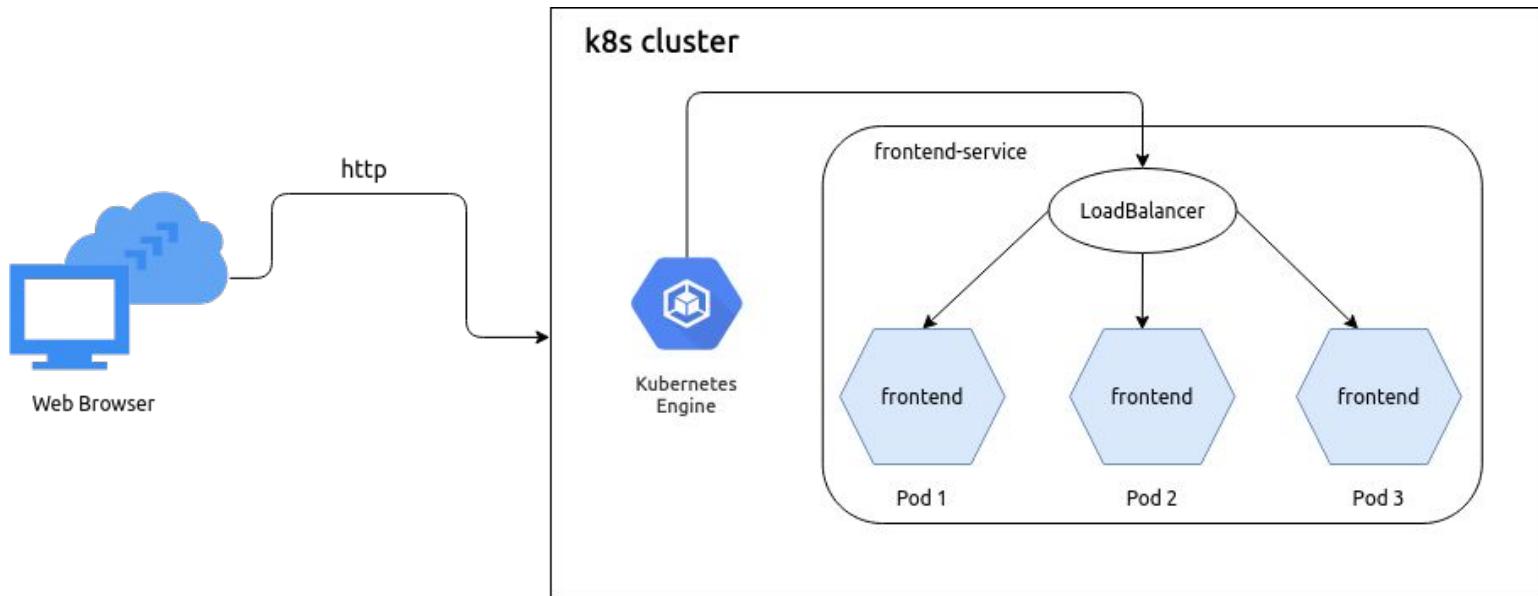
cd ../k8s
gcloud container clusters get-credentials ${CLUSTER_NAME:1:-1}
--region ${REGION:1:-1}
kubectl config set-context --current --namespace=default
kubectl apply -f .

echo "DONE!"
```

# Arquitetura: física



# Arquitetura: k8s



demo



# Repo



[github.com/paulopacitti/titanic](https://github.com/paulopacitti/titanic)

An aerial photograph of a Canadian icebreaker ship sailing on a dark blue ocean. The ship is red with a white superstructure and features the Canadian flag on its mast. It is moving from the bottom left towards the top right, creating a white wake behind it.

**obrigado!**